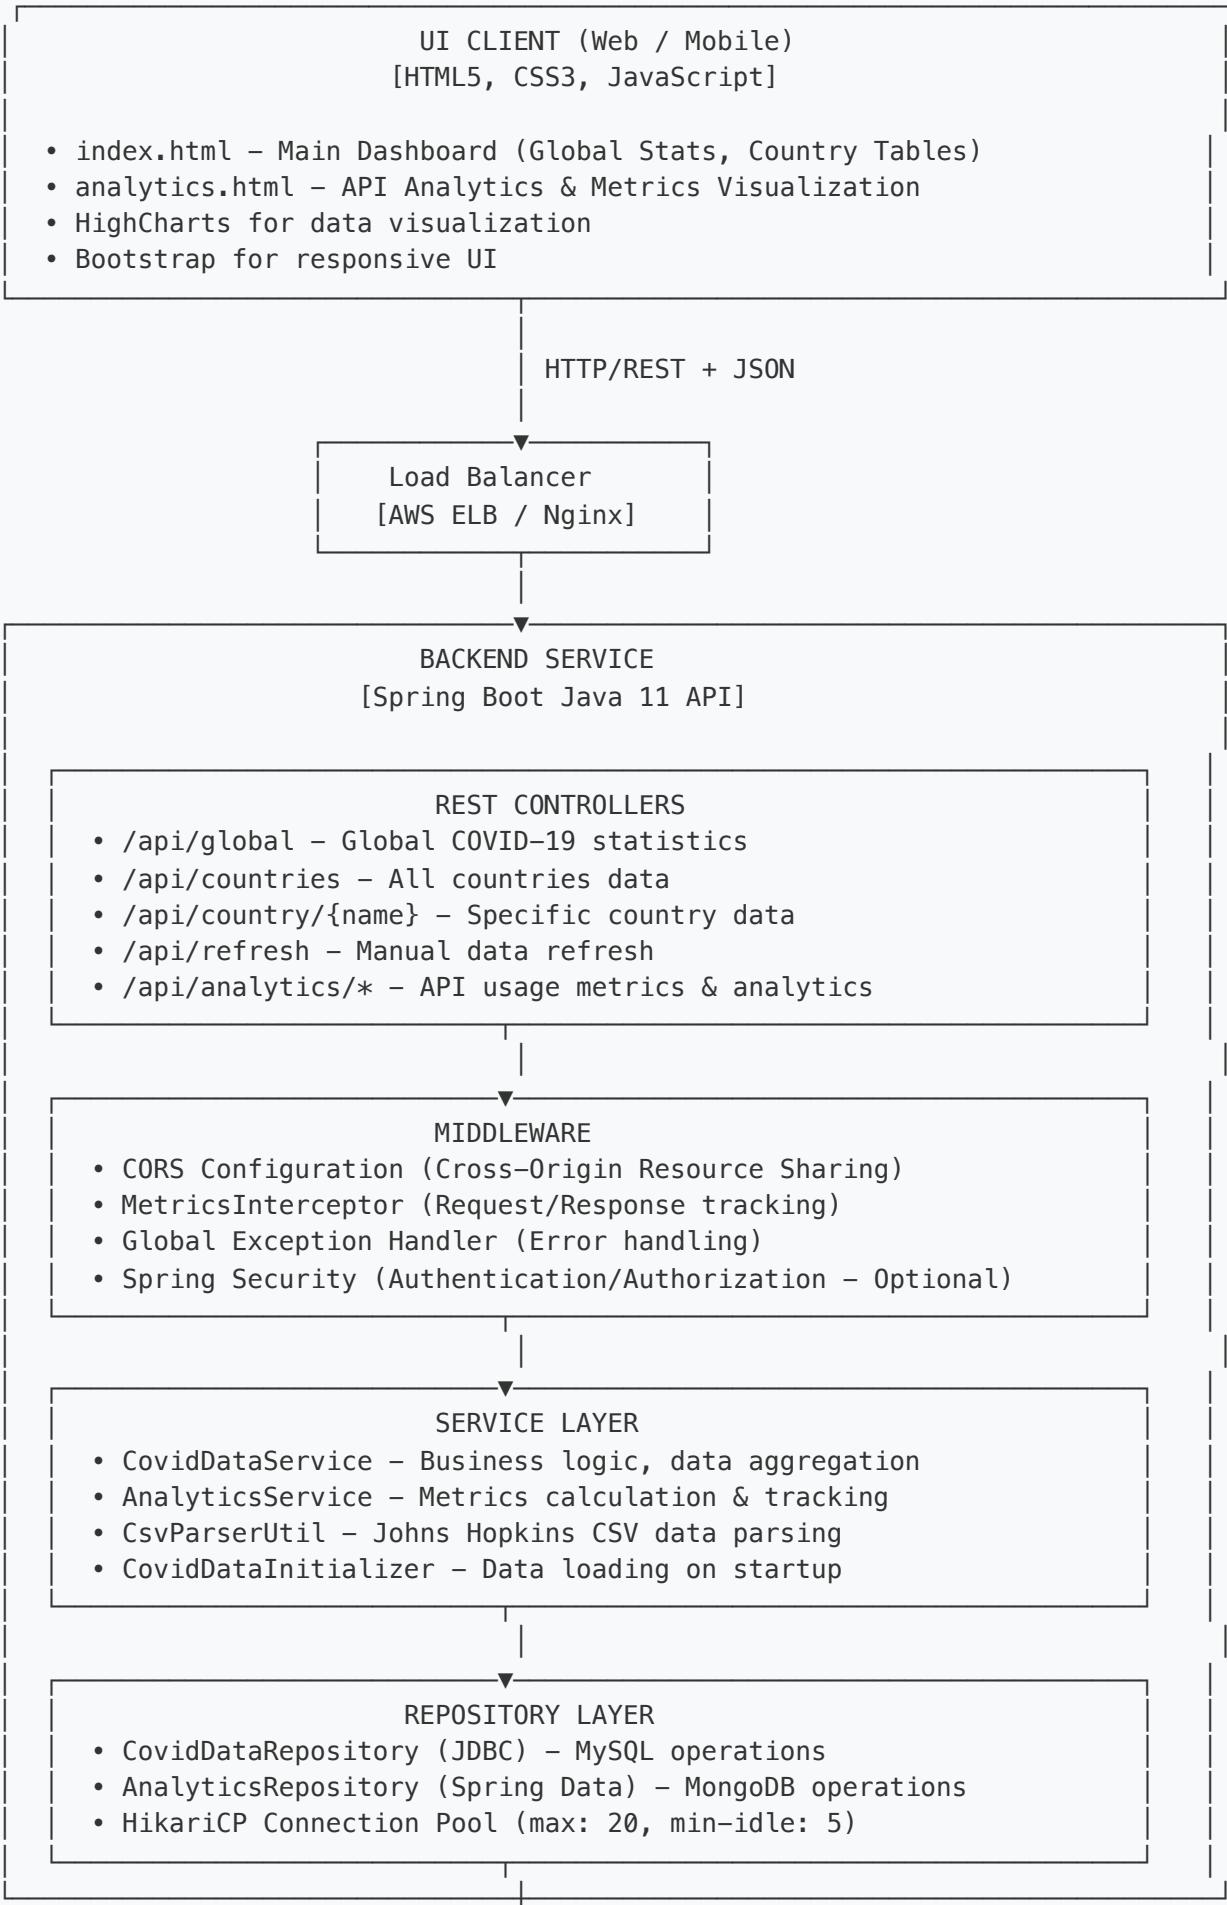


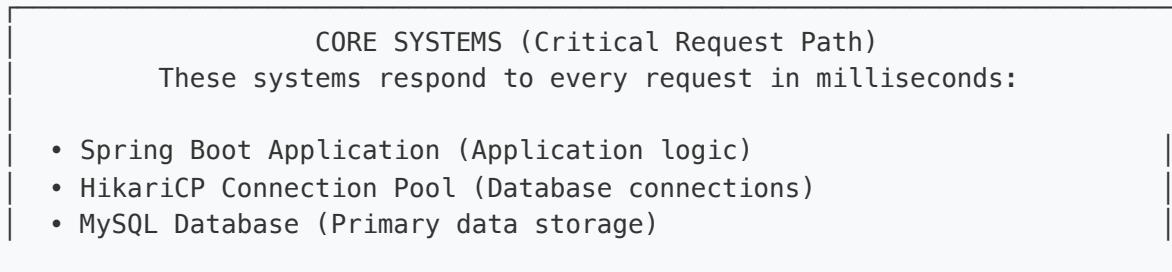
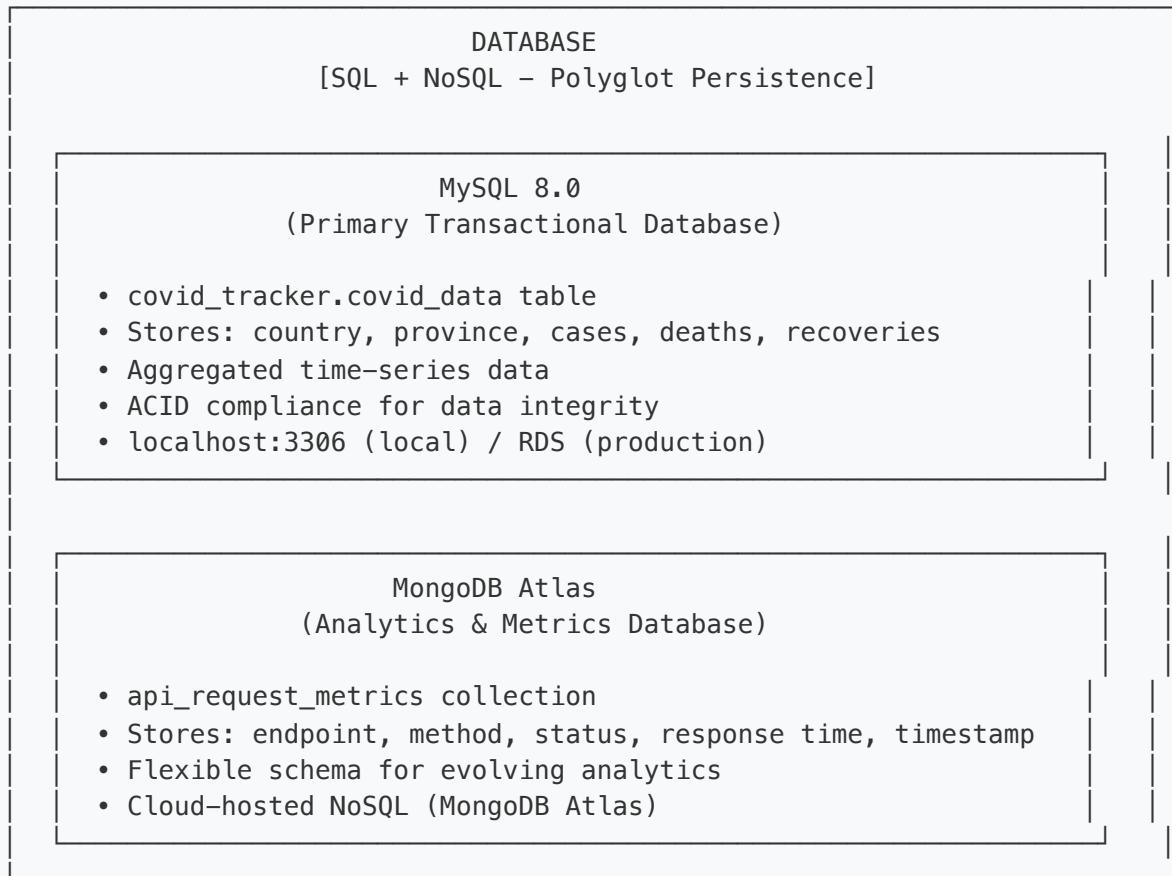
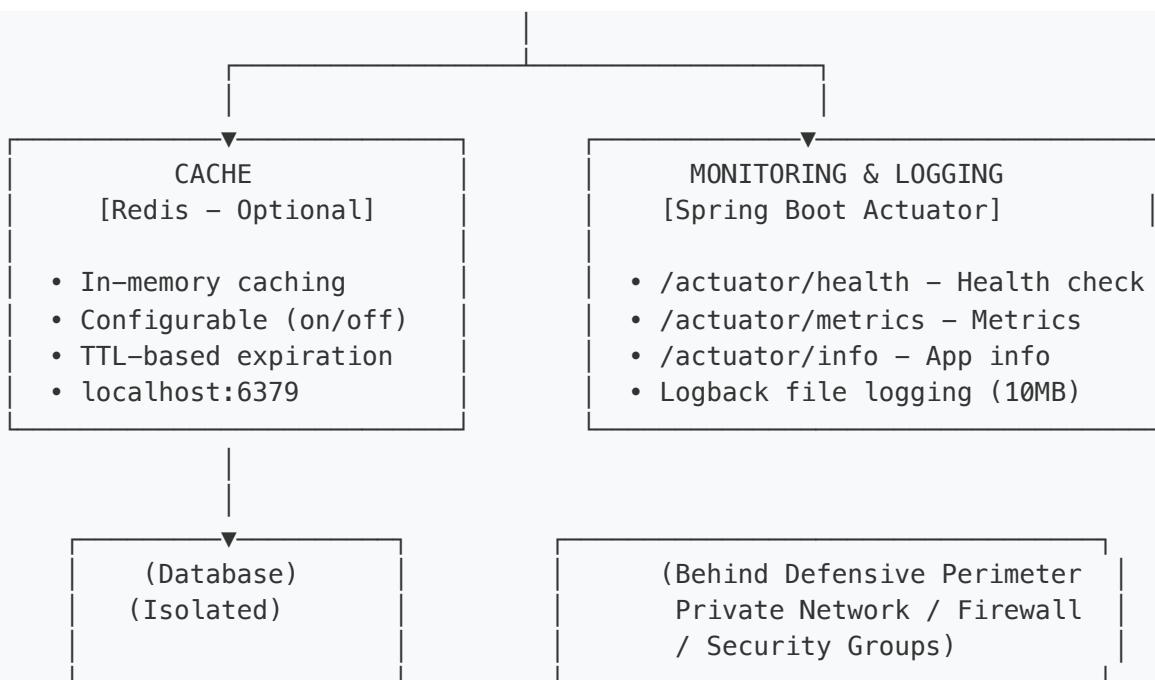
# COVID-19 Tracker - System Architecture Diagram

---

## High-Level System Design

---

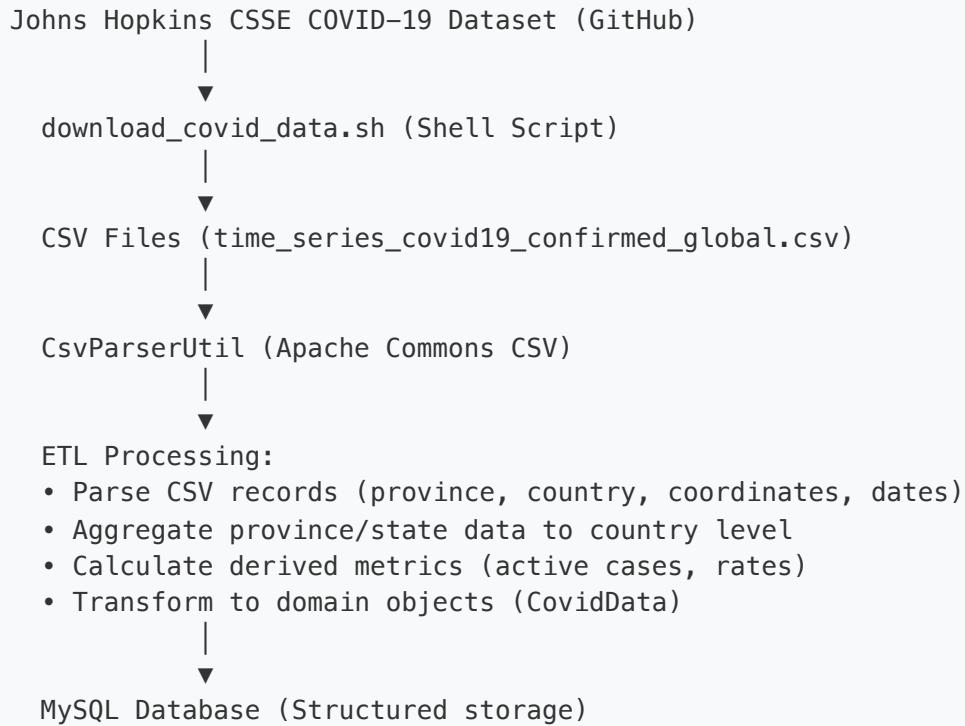




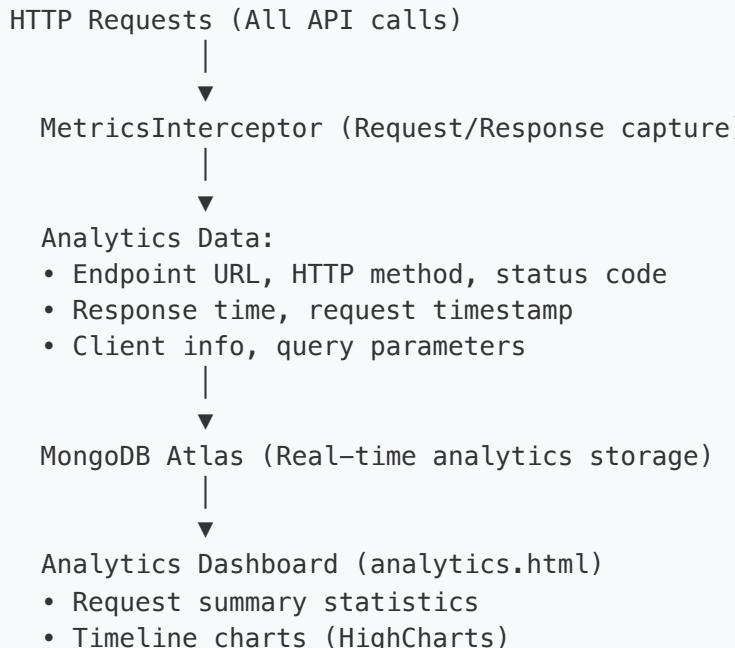
- Redis Cache (Optional – Performance optimization)

**BIG DATA / ANALYTICS PLATFORM**  
**[ETL & Data Processing – Batch & Streaming]**

**DATA INGESTION PIPELINE**



**ANALYTICS PIPELINE**



- Endpoint performance metrics

Receives copies of data from:

- DB changes (via application logs)
- Application logs (Logback rolling file appender)
- Business events (API request metrics)

(Optional insights can be sent back via APIs or to Backend)

### AI / ML ENGINE (Future) [AI / ML + Learning]

- Predictive Analytics (Forecast case trends)
- Anomaly Detection (Unusual data patterns)
- Pattern Recognition (Spread patterns)
- Recommendation Engine (Risk assessments)

Potential ML Models:

- Classical ML: Linear Regression, Decision Trees, Random Forests
- Deep Learning: LSTM/GRU for time-series forecasting
- Clustering: K-means for country grouping by patterns

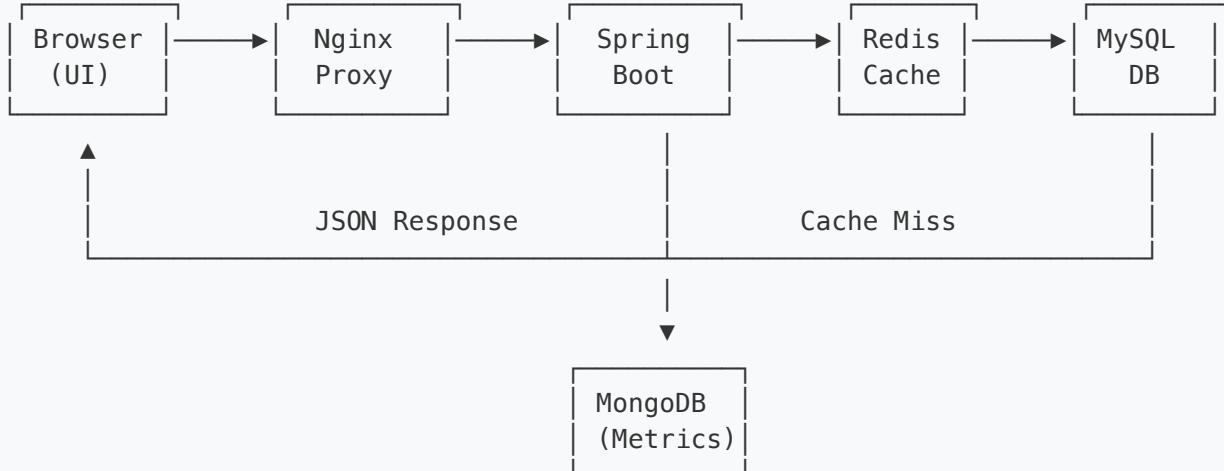
Technologies:

- Python, TensorFlow, PyTorch, scikit-learn
- Jupyter Notebooks for model development
- MLflow for model versioning

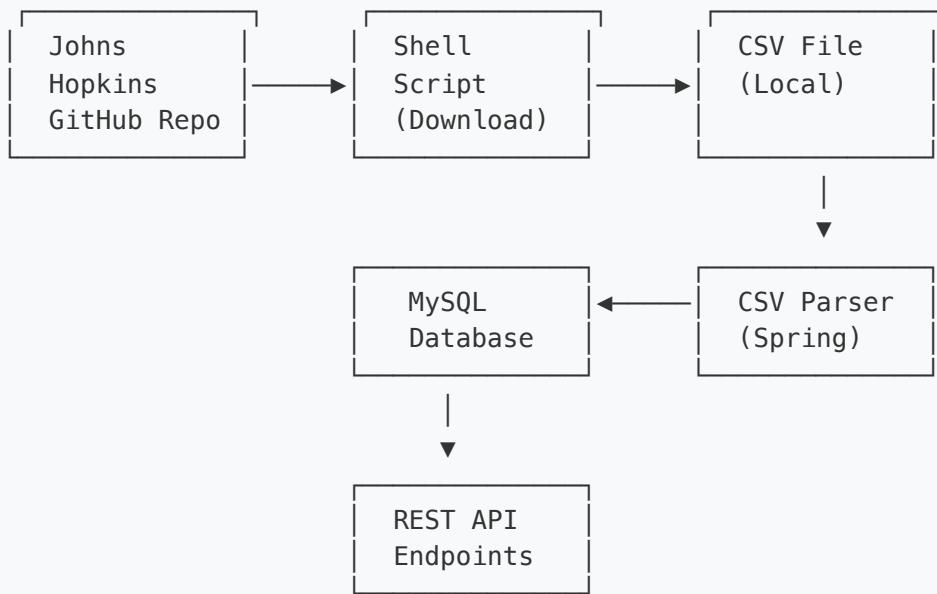
(Optional predictions can be sent back via APIs or Backend)

## Data Flow Diagrams

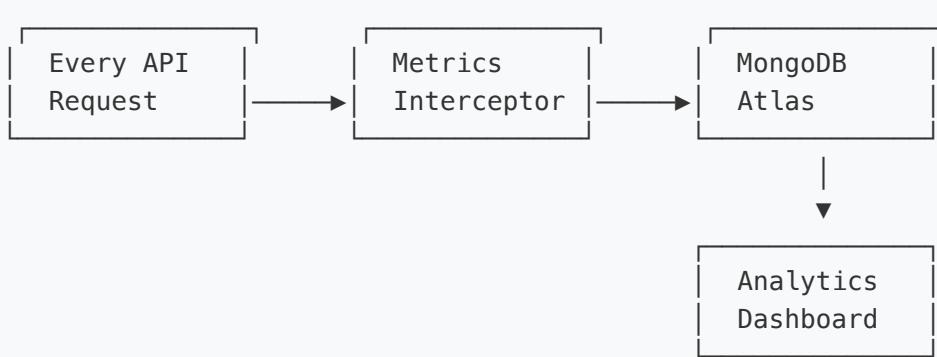
### 1. Primary Data Flow (User Request)



## 2. Data Ingestion Flow (ETL Pipeline)



## 3. Analytics Flow



# Technology Stack by Layer

---

## Frontend Layer

- HTML5 (Semantic markup)
- CSS3 (Responsive design, Bootstrap 5)
- JavaScript (ES6+, Vanilla JS)
- HighCharts.js (Data visualization)
- Fetch API (HTTP requests)

## Application Layer

- Java 11
- Spring Boot 2.7.14
- Spring Web MVC (REST Controllers)
- Spring JDBC (Database access)
- Spring Data MongoDB (Analytics)
- Spring Data Redis (Caching)
- Spring Boot Actuator (Monitoring)
- SpringDoc OpenAPI (Swagger documentation)
- Apache Commons CSV (CSV parsing)
- HikariCP (Connection pooling)
- Logback (Logging)
- Maven (Build & dependency management)

## Data Layer

- MySQL 8.0 (Primary relational database)
  - InnoDB storage engine
  - ACID transactions
  - B-Tree indexes
- MongoDB Atlas (Analytics NoSQL database)
  - Document-based storage
  - Flexible schema
  - Cloud-hosted
- Redis (Optional in-memory cache)
  - Key-value store
  - TTL-based expiration
  - localhost:6379

## Infrastructure Layer

- AWS EC2 (Compute instances)
- Nginx (Reverse proxy, load balancing)
- Systemd (Service management)
- Cron (Scheduled data refresh)
- AWS Elastic IP (Static IP addressing)
- AWS Security Groups (Firewall rules)
- SSH (Secure remote access)

## Key Architecture Patterns

---

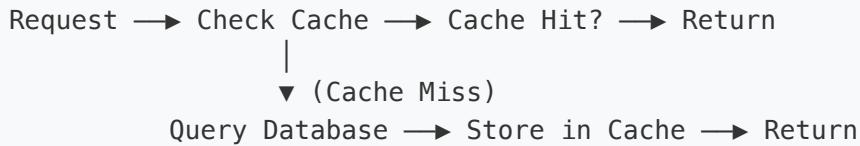
### 1. Layered Architecture (N-Tier)

Controller → Service → Repository → Database  
(REST)        (Business)    (Data Access)    (Storage)

### 2. Polyglot Persistence

MySQL (Transactional) + MongoDB (Analytics) + Redis (Cache)

### 3. Caching Strategy



## 4. MVC Pattern (Frontend)

Model (Data) → View (HTML) → Controller (JavaScript)

## 5. Repository Pattern (Backend)

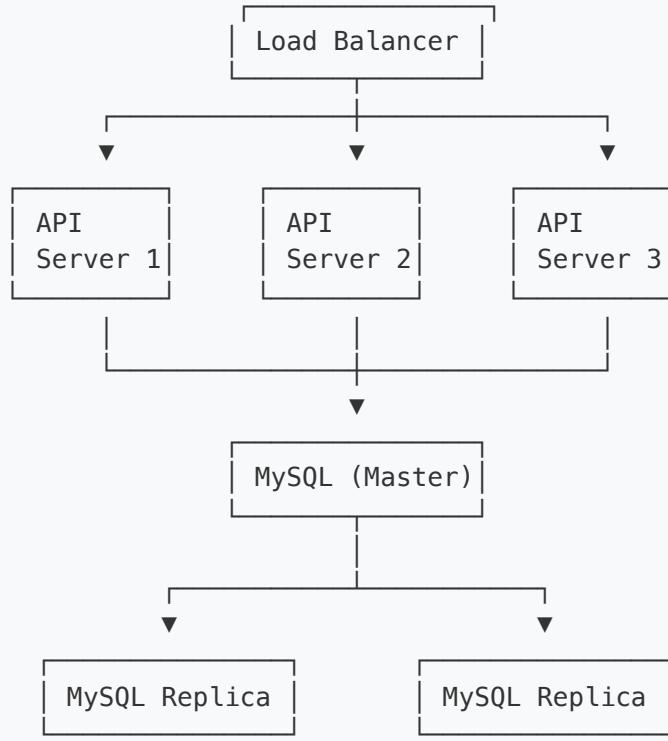
Service → Repository Interface → Repository Implementation → Database

## Security Considerations

- CORS Configuration (Restricted origins)
- AWS Security Groups (Port restrictions)
- Environment-based configuration (secrets)
- Connection pooling limits (DoS prevention)
- Actuator endpoint security
- Database credentials management
- HTTPS/SSL (Production – via AWS ELB/Nginx)

## Scalability & Performance

### Horizontal Scaling



## Caching Strategy

- **Redis Cache:** Optional in-memory caching for frequently accessed data
- **HikariCP Pool:** Connection reuse (5-20 connections)
- **HTTP Caching:** Browser-level caching headers

## Database Optimization

- **Indexes:** On country, province, date columns
- **Connection Pooling:** HikariCP with configurable limits
- **Query Optimization:** Efficient aggregation queries

## Monitoring & Observability

---

**Spring Boot Actuator Endpoints:**

- /actuator/health – Application health status
- /actuator/metrics – JVM & application metrics
- /actuator/info – Application metadata

**Logging:**

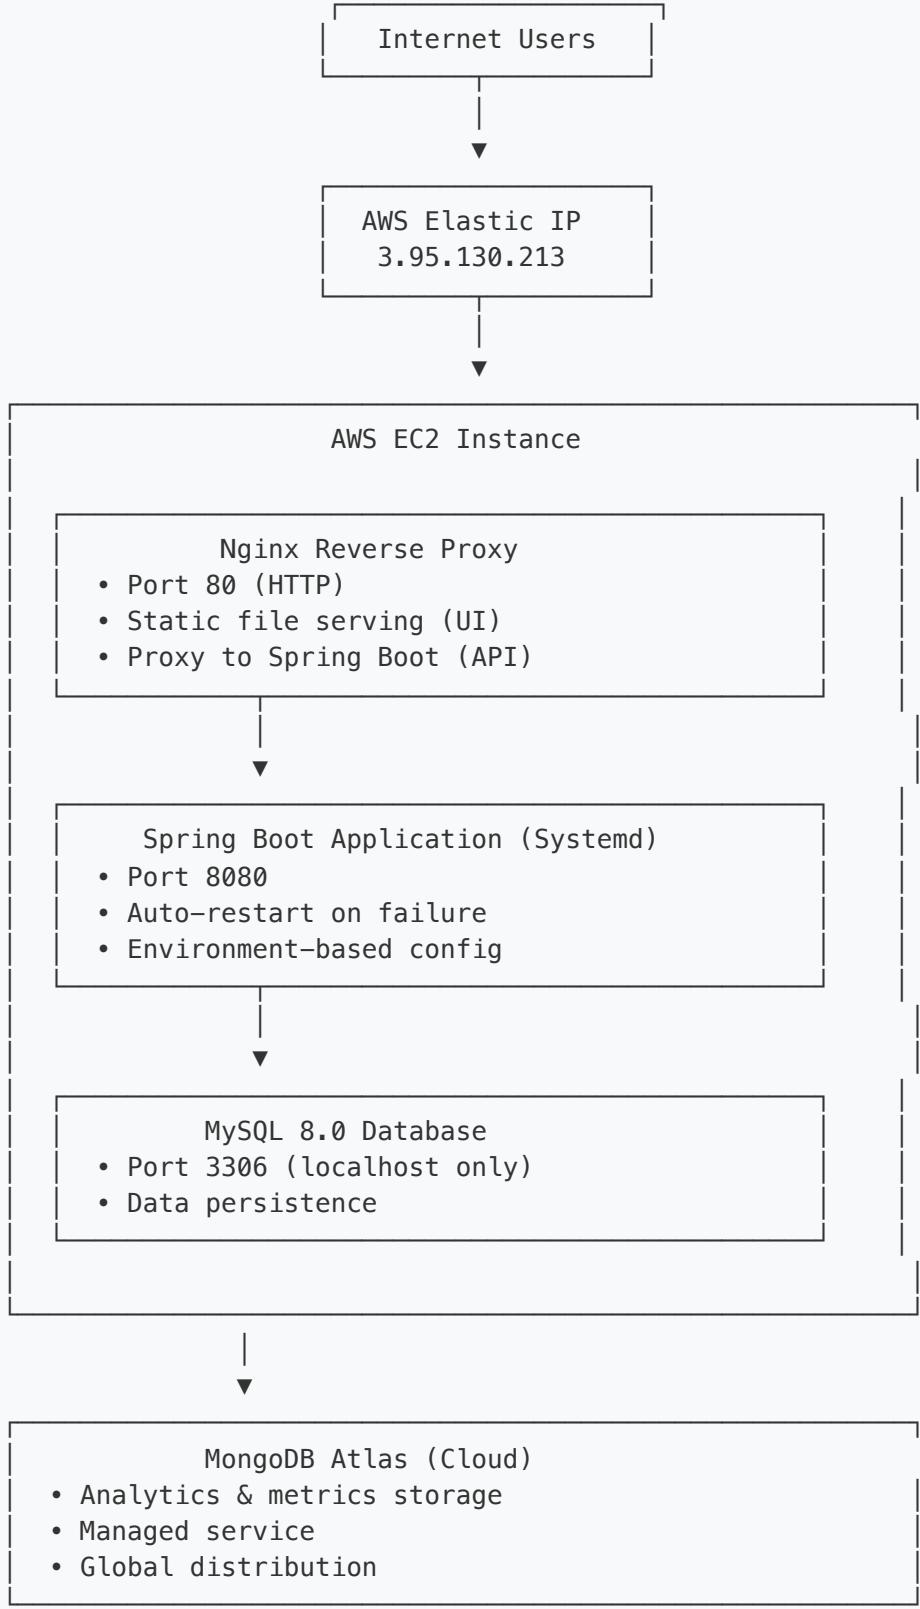
- Logback with rolling file appender
- 10MB file size limit
- 10 files retention
- Gzip compression for archived logs

**Analytics Dashboard:**

- Real-time API usage metrics
- Request/response time tracking
- Endpoint performance visualization

## Deployment Architecture (AWS)

---



## API Endpoints Summary

### COVID Data Endpoints

- `GET /api/global` - Global statistics
- `GET /api/countries` - All countries data
- `GET /api/country/{name}` - Specific country data
- `POST /api/refresh` - Manual data refresh

## Analytics Endpoints

- `GET /api/analytics/summary` - Request summary statistics
- `GET /api/analytics/timeline` - Timeline metrics

## Monitoring Endpoints

- `GET /actuator/health` - Health check
- `GET /actuator/metrics` - Application metrics
- `GET /actuator/info` - Application info

## Documentation

- `GET /swagger-ui/index.html` - Interactive API documentation

---

**Created for:** COVID-19 Tracker Full Stack Application

**Version:** 1.0.0

**Last Updated:** January 2026